


Mosaics





PRESENTED BY

Brad Aimone



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia LLC, a wholly owned subsidiary of Honeywell International Inc. for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.





Artificial neural networks

- Generic layers of non-linear nodes
- SGD optimization of weights
- Powerful machine learning capabilities through learning sequential non-linear mappings and function approximation


Spiking neural algorithms


- Hand-crafted circuits of spiking neurons
- Model of parallel computation
- Energy efficiency through event-driven communication and high fan-in logic



Neuroscience-constrained algorithms


- Circuit architecture based on local and regional neural connectivity
- Computation incorporates broad range of neural plasticity and dynamics
- *Generally still unexplored from algorithms perspective*






Neural Networks

- Better-suited for neuromorphic hardware than many other machine learning techniques
- ANNs only became broadly world class when they reached substantially large sizes




Matrix Multiplication

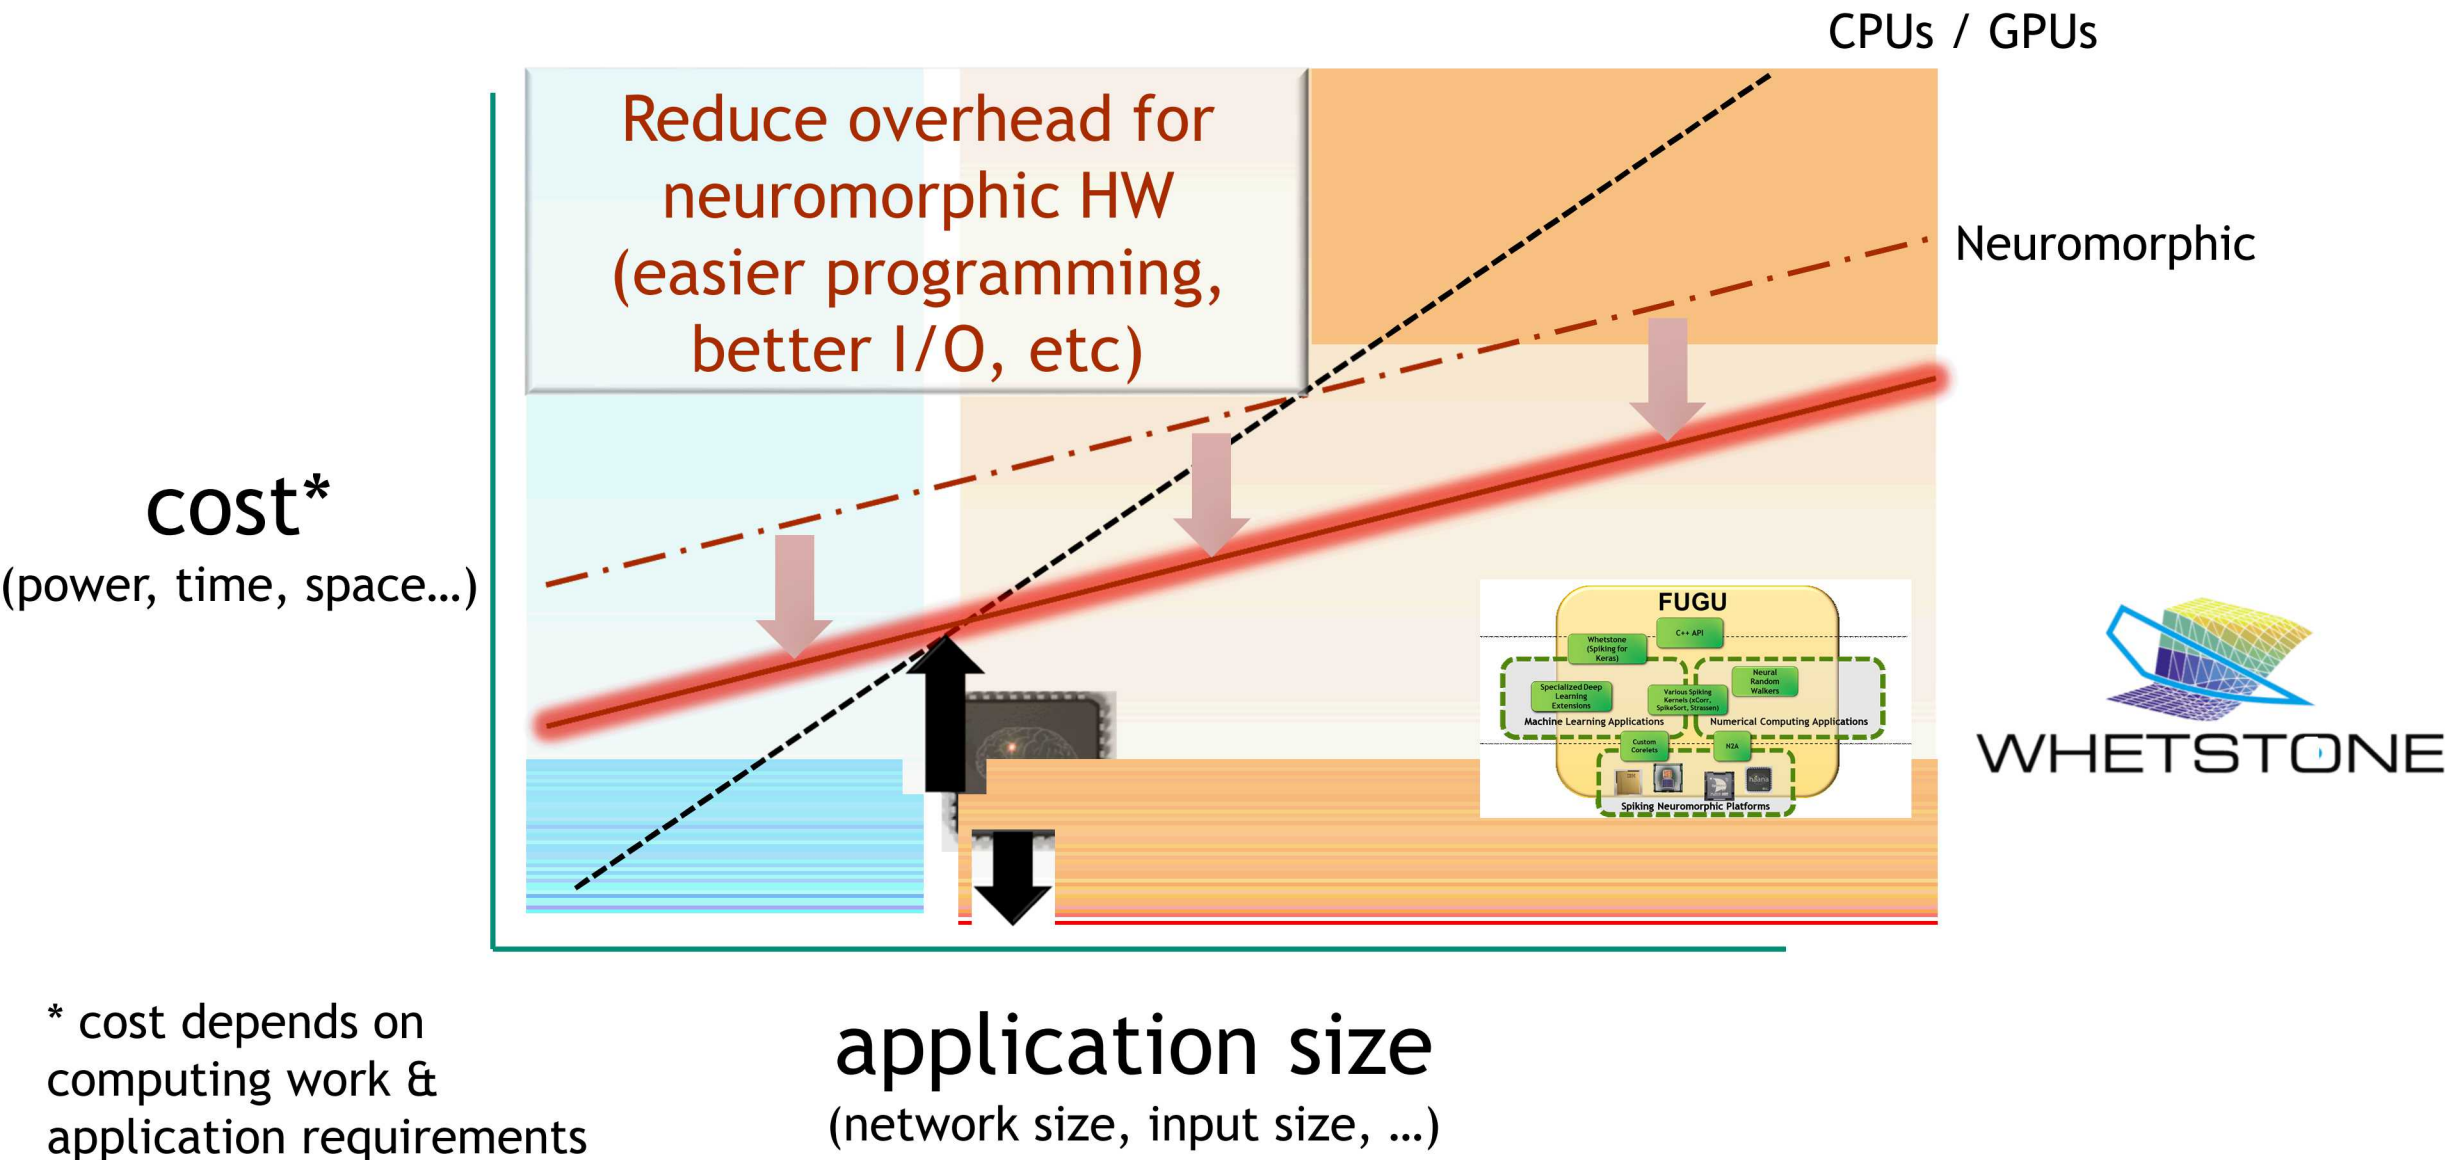
- Neural algorithms can improve implementation of Strassen-techniques
- Strassen techniques only make sense for large matrix multiplications




Partial Differential Equations


- Neural algorithms can efficiently implement Monte Carlo solutions for solving diffusion-based PDEs
- Monte Carlo methods make most sense for high-dimensional PDEs







* cost depends on
computing work &
application requirements







Memory is cheap. We should take advantage of that.

More synapses (whether as arrays or tables or whatever) is not really the problem


Communication is expensive. And it scales poorly with neurons and synapses...

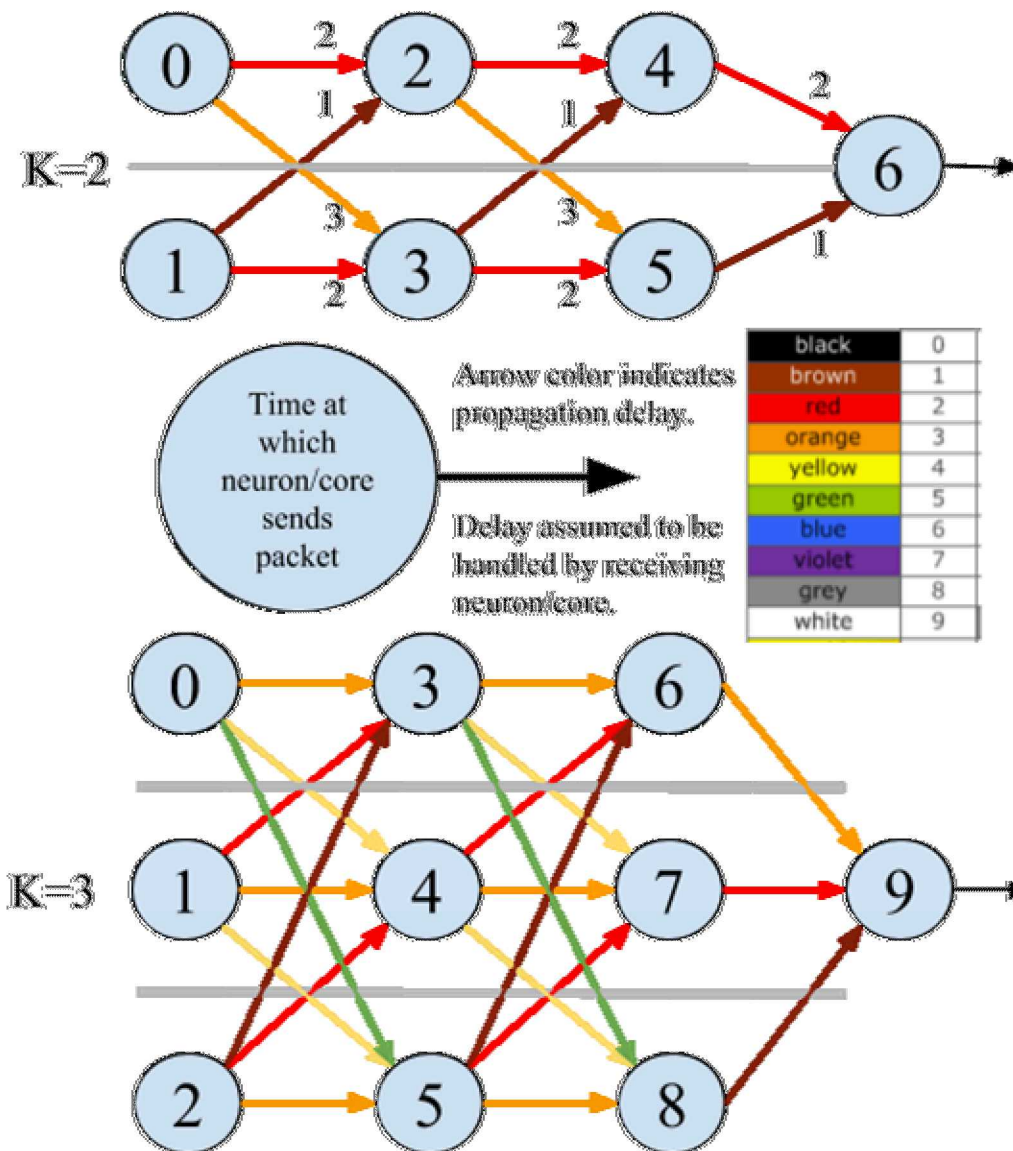
Can we formulate an approach to mitigate this?

Neurons are not as cheap. And we need a lot of them.

But depending on the algorithm, they are often unused for large periods of time.

Can we better take advantage of what we have?






Divide each layer into K temporal-groups. Assign each group a relative-firing-time (accounting for the global minimum delay). Create a data structure that allows looking up the relative-firing-time based on the neuron-id.


When making a connection from one neuron to another, set the delay to be the relative-firing-time of the postsynaptic-neuron minus the relative-firing-time of the presynaptic-neuron.


When scheduling the spikes of a spike-source array, schedule them at times offset from the sample presentation time based on the relative-firing-time of the source-neuron.


All the neurons for the output layer of a classifier should be placed in the same temporal-group, which simplifies reading output.

In the case of a semantic segmenter (for example), the output would also have to be temporally-staggered, and so some decoding would be necessary.











Obviously this won't work for everything

But it will work for a lot of potential applications




Machine Learning

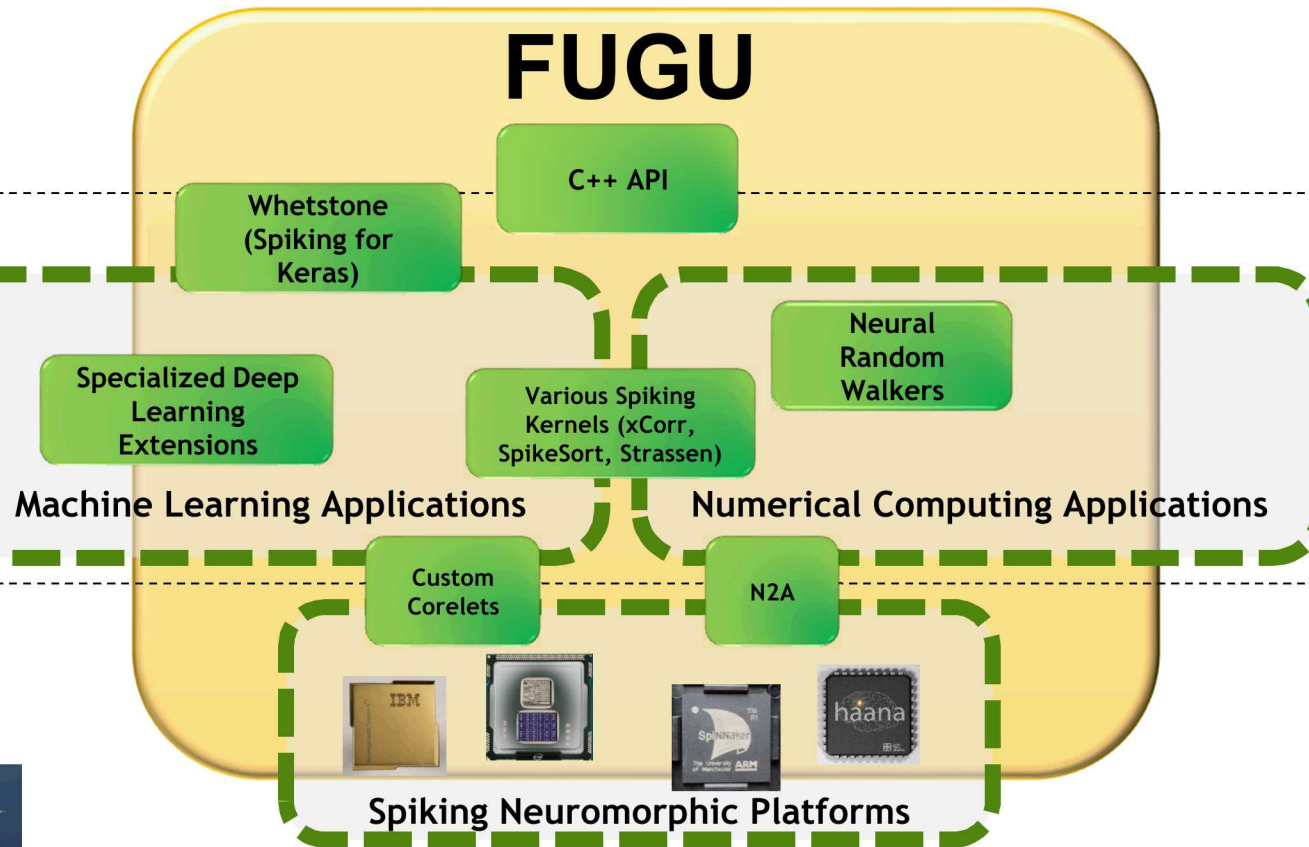
Whetstone

Convolutions

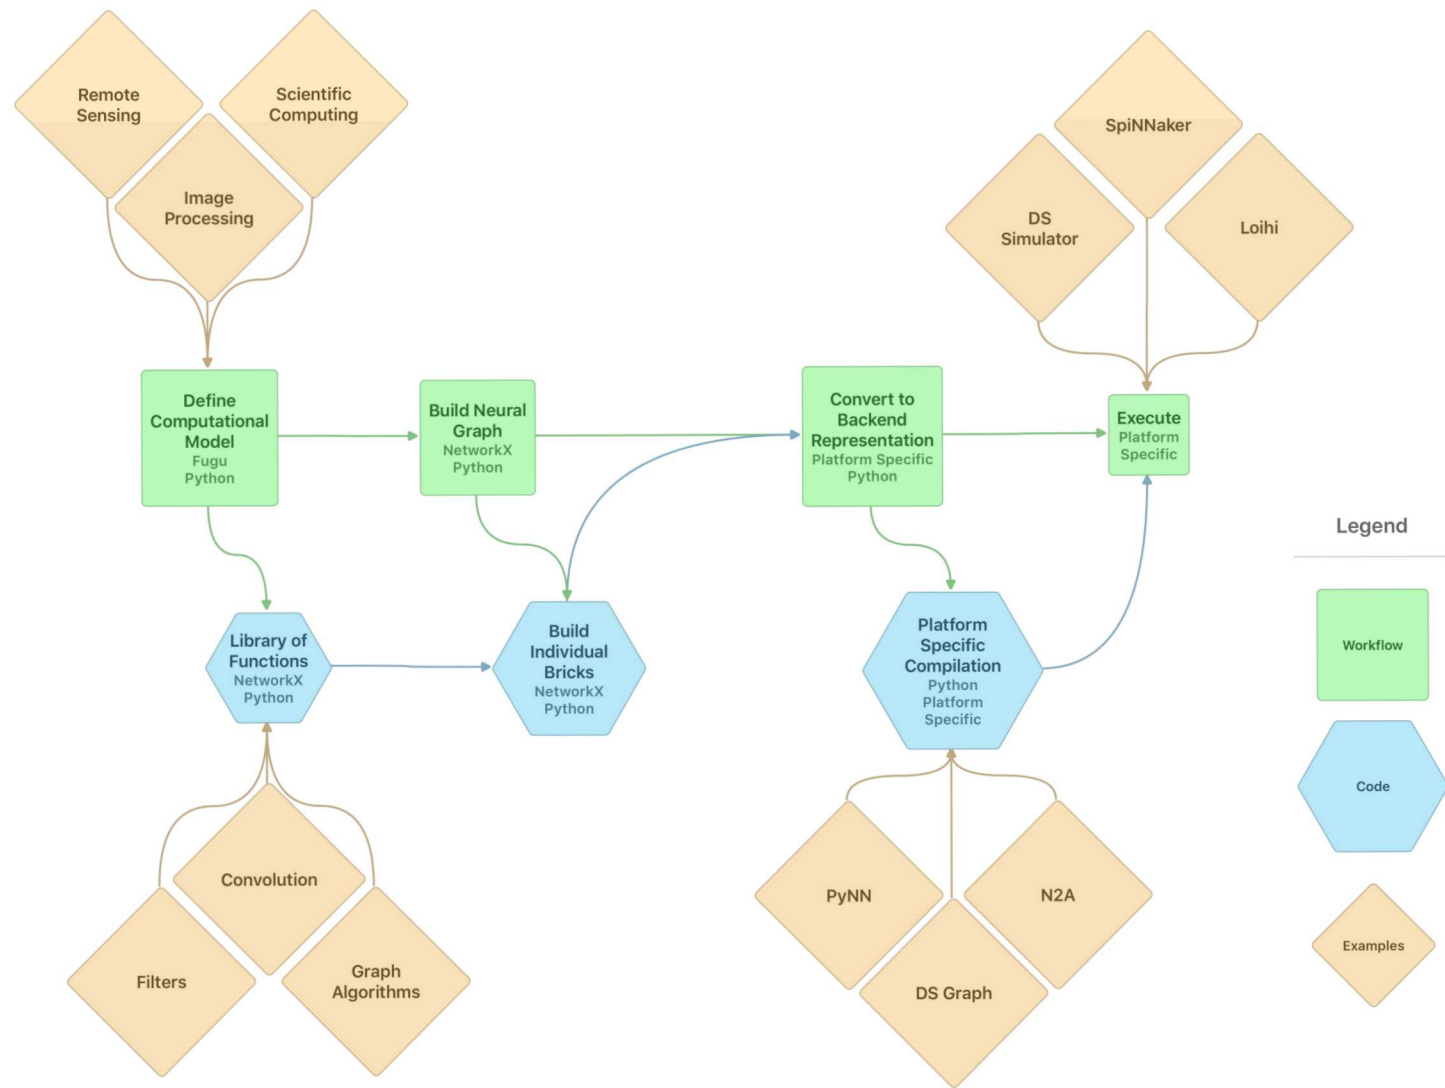
k-Nearest Neighbor

Support Vector Machines





Fugu = pufferfish (why? Pufferfish have spikes...)



Thanks everyone for coming to NICE 2019!